



AlBuster

Under Supervision of:

Dr.Ahmed Madbouly

AlBuster

AI-Generated Image Detection Research

Name	ID	Department	Grad		
			Discussion (60)	Follow-up (20)	Presentation (20)
Moustapha Muhamed Hassaan	352251015	Statistics and Computer Science			
Moustapha Emad Muhamed	352251013	Statistics and Computer Science			
Moustapha Hossam Abdel-Rahim	352251007	Statistics and Computer Science			
Muhamed Wael Kamal	352250977	Statistics and Computer Science			
Muhamed Sayed Abdel-Samee	352250947	Statistics and Computer Science			
Nada Bakr Nagah	352251041	Statistics and Computer Science			

JUDGMENT COMMITTEE

Examination & Evaluation Panel — AlBuster Research Project

#	Name	Signature
1		
2		
3		
4		
5		

APPROVED BY

Department of Statistics and Computer Science

Table of Figures

Figure	Page	Description
Figure 1	29	High-level system architecture showing the parallel analysis pipeline
Figure 2	33	FFT comparison — (a) Real photograph showing smooth radial decay, (b) AI-generated image showing grid artifacts in the frequency domain
Figure 3	34	ELA comparison — (a) Real photograph showing uniform low error levels, (b) AI-generated image showing non-uniform glowing regions indicating synthetic content
Figure 4	37	Noise analysis comparison — (a) Real photo showing bright textured noise residual, (b) AI image showing dark patterned noise residual
Figure 5	41	ViT architecture for AI image detection — the image is split into patches, processed through 12 transformer encoder layers, and classified as real or fake
Figure 6	43	Transfer learning workflow — pre-trained knowledge is adapted for the specific classification task
Figure 7	45	Screenshot of the training output showing epoch-by-epoch metrics
Figure 8	48	Adaptive ensemble decision flow — threshold adjusts based on model agreement
Figure 9	50	Multi-scale inference pipeline — four views are analyzed independently and averaged

Table of Abbreviations

Abbreviation	Full Name
AI	Artificial Intelligence
AAAI	Association for the Advancement of Artificial Intelligence
GPT	Generative Pre-trained Transformer
CNN	Convolutional Neural Network
ResNet	Residual Network
VGGNet	Visual Geometry Group Network
DenseNet	Densely Connected Convolutional Networks
RELU	Rectified Linear Unit
ROC-AUC	Receiver Operating Characteristic – Area Under the Curve
ViT	Vision Transformers
GAN	Generative Adversarial Networks
LLM	Large Language Models
XceptionNet	Extreme Inception
DFT	Discrete Fourier Transform
DWT	Discrete Wavelet Transform
VLMs	Vision Language Models
MLLMs	Multimodal Large Language Models

Table of Contents

Table of Figures	i
Table of Abbreviations	ii
Abstract	v
Chapter 1: Introduction	1
1.1 Background and Motivation	2
1.2 Proposed Methodology: A Dual-Model Approach	2
1.3 Research Objectives	3
Chapter 2: Related Work	4
2.1 Detecting AI-Generated Images Using a Hybrid ResNet-SE Attention Model	5
2.2 C2PA	7
2.3 Google SynthID	10
2.4 Methods and Trends in Detecting AI-Generated Images	11
2.5 Summary and Research Gap	13
Chapter 3: Methodology and Implementation	14
3.1 The First Model — Vision Transformer (ViT)	15
3.2 The Second Model — Multi-Engine Ensemble System	31
3.3 Deployment and Interface Design	56
Chapter 4: Comparative Analysis of AI Image Detection Models	65
4.1 Model 1 — The Standalone Vision Transformer	66
4.2 Model 2 — The Multi-Engine Ensemble Detection System	67
4.3 Side-by-Side Comparison	68
4.4 Discussion	71

Chapter 5: Conclusion and Future Work	72
5.1 Summary of Work	73
5.2 Key Findings	74
5.3 Limitations	75
5.4 Future Work	76
5.5 Closing Remarks	77
Chapter 6 :Diagrams	79
References	83

Abstract

The rapid advancement of generative AI has made synthetic images increasingly indistinguishable from real photographs, enabling misinformation, fraud, and digital manipulation at an unprecedented scale. This paper presents AlBuster, a dual-model AI-generated image detection system designed to address this challenge through two complementary architectural approaches. The first model is a fine-tuned Vision Transformer (ViT), pre-trained on ImageNet-21k and specialized on AI detection datasets, achieving 94.2% accuracy with ~45ms inference time. The second is a multi-engine ensemble system that combines a fine-tuned ViT with deep learning backup models, physics-based noise forensics, FFT frequency analysis, and Error Level Analysis, achieving 99.40% accuracy and 100% recall on its test set.

Both models are deployed through a web interface built on Gradio and FastAPI, and integrated into an Android mobile application using an MVVM architecture with Kotlin and Ktor. The system is publicly accessible via Hugging Face Spaces.

The results demonstrate that combining transformer-based global pattern recognition with physics-based signal forensics produces a detection framework that is robust, interpretable, and practically deployable — laying groundwork for next-generation AI authenticity verification tools.

C H A P T E R 1

Introduction

1.1 Background and Motivation

Artificial intelligence art generators are trained on billions of existing images. When a user enters a prompt, the AI art generator builds an image by combining aspects of its training data into a single composite output. Meanwhile, deepfakes are trained on photographs and videos of one subject to replicate that subject, and can depict a person saying or doing something they never did. Both technologies can be used harmlessly or maliciously.

Scammers use deepfake technology to impersonate celebrities or public figures, creating endorsements for fake products and stealing financial information. Deepfakes can portray government officials making statements they never made. They also serve as instruments for blackmail, producing near-indistinguishable fake images to extort money or favors.

Beyond deepfakes, cybercriminals use AI-generated art in online romance scams, fraudulent donation campaigns, and "Bait and Vanish" scams where buyers pay for products that never arrive. The widespread availability of generative AI has raised serious concerns among researchers and policymakers, making robust detection methods critically important.

1.2 Proposed Methodology:

To address these challenges, this research proposes the development and evaluation of two separate and distinct models designed to capture both local and global generative signatures:

Model 1 — The Fine-Tuned Vision Transformer (ViT): The first model applies the Transformer architecture to image classification by dividing each image into fixed-size patches and processing them through multi-head self-attention layers. Pre-trained on ImageNet-21k and fine-tuned on specialized AI detection datasets, this model captures global structural inconsistencies and generative fingerprints that traditional CNN-based detectors typically miss. A carefully tuned classification threshold further improves sensitivity to subtle generative signals.

Model 2 — The Multi-Engine Ensemble System: The second model takes a broader forensic approach, combining multiple independent analysis engines into a single weighted verdict. A fine-tuned ViT serves as the primary classifier, supported by backup deep learning models and a physics-based noise forensics engine that analyzes statistical properties of image noise without relying on any training data. Visual forensic tools — including FFT frequency analysis and Error Level Analysis — provide additional human-interpretable evidence alongside the automated verdict.

1.3 The main objectives of this research:

- To design and evaluate a fine-tuned Vision Transformer (ViT) capable of detecting AI-generated images by capturing global structural inconsistencies and generative fingerprints that traditional CNN-based architectures typically miss.
- To develop a multi-engine ensemble detection system that combines deep learning classification with physics-based forensic analysis — including FFT frequency analysis, Error Level Analysis, and noise pattern forensics — to improve robustness and interpretability.
- To conduct a comparative analysis of both models across key dimensions including accuracy, inference speed, forensic transparency, and robustness, in order to evaluate the trade-offs between a focused single-model approach and a multi-layered ensemble strategy.
- To deploy both models in accessible, real-world interfaces: a web application via Hugging Face Spaces and an Android mobile application, making AI detection available beyond research environments.

C H A P T E R 2

Related Work

This chapter reviews the existing literature most relevant to AlBuster across three dimensions: model-based detection approaches, provenance and watermarking standards, and broader trends in the field. Together, these works inform the architectural decisions and design goals of the proposed system. The works were selected to represent the two dominant paradigms in the field — detection-based and provenance-based approaches — and to identify the limitations that motivated AlBuster's dual-model design.

2.1 Detecting AI-Generated Images Using a Hybrid ResNet-SE Attention Model

The paper explores the detection of AI-generated images by evaluating three established baseline architectures and proposing a high-performance hybrid model, SE-ResNet50, which integrates attention mechanisms to identify subtle synthetic artifacts.

2.1.1 Standard Baseline Models

The study first benchmarks three prominent CNN architectures to establish a performance baseline:

Model	Architecture	Accuracy	Notes
ResNet-50	Residual learning with skip connections	93.46%	May struggle with very fine-grained real vs. synthetic differences
VGGNet (VGG-16)	Stack of uniform 3×3 conv layers	94.06%	Strong generalization; high memory and computational cost
DenseNet (DenseNet-)	Densely connected layers; max feature	95.98%	Competitive accuracy; very high

Model	Architecture	Accuracy	Notes
121)	reuse		computational cost

2.1.2The Proposed Hybrid Model: SE-ResNet50

The core contribution of the research is the SE-ResNet50 model, which augments the ResNet-50 backbone by integrating Squeeze-and-Excitation attention blocks into every residual block. This hybrid approach dynamically recalibrates the importance of feature channels, allowing the network to focus on informative generative cues while suppressing noise. The SE mechanism operates through three stages: Squeeze (Global Average Pooling to create a channel descriptor), Excitation (two fully connected layers with ReLU and sigmoid activations to learn channel attention scores), and Scaling (weighting the original feature maps by the attention scores).

2.1.3Experimental Setup and Training

The models were trained and evaluated on the CIFAKE dataset, which contains 50,000 images split evenly between real and AI-generated categories. All images were resized to 224×224 pixels. Standardized hyperparameters were used across all models for fair comparison: Adam optimizer, learning rate of 0.0001, Cross-Entropy Loss, batch size of 32, and 10 training epochs.

2.1.4Performance Results

The SE-ResNet50 model emerged as the superior architecture, achieving a test accuracy of 96.12% and a ROC-AUC score of 0.9862. It outperformed both traditional CNN baselines and the transformer-based ViT. The paper attributes this success to the model's ability to selectively emphasize low-level artifacts and spatial distortions often introduced by GANs and diffusion networks.

2.1.5Critical Analysis

This paper demonstrates that attention mechanisms can meaningfully improve detection by helping the model focus on the subtle artifacts that generative models leave behind. The strong results on CIFAKE suggest that hybrid CNN architectures are a viable direction. However, the reliance on a single dataset and a binary real/fake classification leaves questions about how well this approach generalizes to

newer generative models like diffusion-based systems.

2.2 C2PA

2.2.1 What C2PA Is

C2PA — the Coalition for Content Provenance and Authenticity — is a Joint Development Foundation project that publishes royalty-free, open technical specifications for attaching verifiable provenance metadata to digital media. A C2PA-enabled camera, photo editor, or AI platform can embed a signed record inside any image, video, audio file, or document. That record travels with the file and can be verified by anyone using any compliant tool, with no internet connection or central database required. The core data structure is called a C2PA Manifest (also marketed as a Content Credential), containing assertions about who created the content, when, what tools were used, whether AI was involved, and every meaningful edit made since capture. The manifest is digitally signed — any tampering breaks the signature and is immediately detectable.

2.2.2 The Problem C2PA Solves

Generative AI has made synthetic media indistinguishable from authentic media at a visual level. Deepfake incidents tracked globally surged from approximately 500,000 cases in 2023 to over 8 million in 2025 — a 900% increase in two years. Synthetic content is projected to account for up to 90% of online media by 2026. Detection-only approaches are a losing battle because generative models improve continuously. C2PA takes a different approach: instead of detecting fakes after the fact, it focuses on proving authenticity at the point of creation. The U.S. CISA endorsed content credentials in a January 2025 advisory recommending their adoption by government agencies.

2.2.3 How C2PA Works: Step by Step

Stage	Who Acts	What Happens
1. Signing	Camera, software, or AI platform	Assembles assertions into a claim and signs it with a private key issued by a trusted Certificate Authority
2. Embedding	Signing tool	Stores the signed manifest inside the file (JUMBF container) or as a sidecar file; a hard binding cryptographic hash links manifest to pixel data
3. Verification	Any C2PA-compliant tool or viewer	Reads the manifest, validates the certificate chain, checks the hard binding hash against the current file bytes, and reports pass or fail

2.2.4Who Supports C2PA

Organization	Sector	Role / Implementation
Adobe	Software / Creative tools	Founding member; C2PA signing in Photoshop, Lightroom, Firefly, and Stock
Microsoft	Software / Cloud	Founding member; Project Origin lead; Azure content integrity APIs
BBC	News media	Founding member; Project Origin lead; news content verification workflows
Sony / Leica / Nikon	Camera hardware	C2PA signing built into professional camera bodies
Google	Mobile / AI	C2PA credentials planned for Pixel 9 camera and Imagen AI generator
Stability AI / OpenAI	Generative AI	Content credentials on AI-generated images to disclose synthetic origin

2.2.5Which Cameras and AI Tools Support C2PA

As of early 2026, major implementations include: Leica M11-P (October 2023) — the first consumer camera with built-in C2PA; Sony α9 III / α1 II (2024) — opt-in C2PA support; Samsung Galaxy S25 (January 2025) — credentials on AI-edited photos; Google Pixel 10 (September 2025) — every photo signed by default with hardware-backed keys; and AI image generators including Adobe Firefly, OpenAI DALL-E 3, Sora, and Google Imagen. Midjourney does not currently implement C2PA.

2.2.6 Limitations and Open Challenges

C2PA is powerful but not a complete solution. Key limitations include: strip attacks (a non-C2PA tool can silently remove all credentials by re-saving the file); first-mile trust (C2PA proves a file was signed by a specific device but cannot verify what the camera was pointed at); Trust List maturity (the Certificate Trust List is still maturing); and adoption gaps (most consumer smartphone cameras do not yet sign natively).

2.2.6 Critical Analysis

C2PA shifts the detection problem entirely — rather than analyzing pixels, it establishes trust at the point of creation. This is a fundamentally more scalable approach, but its effectiveness depends entirely on adoption. Since most content online lacks C2PA signatures, detection-based systems remain essential for handling the vast majority of unverified media in circulation.

2.3 Google SynthID

SynthID represents a significant advancement in digital steganography and AI provenance. Rather than relying on fragile metadata or visible watermarks, SynthID utilizes deep learning to embed signals directly into the model's output distribution.

2.3.1 Visual Content (Images and Video)

For images and video, SynthID employs a dual-network architecture of an encoder and a detector. The encoder makes subtle, pixel-level adjustments designed to be perceptually invisible to humans, often operating within the frequency domain or latent space of the image. The watermark is engineered to survive lossy transformations such as JPEG compression, cropping, and color jittering by ensuring the embedded pattern is spatially distributed across the entire frame.

For audio, SynthID converts the sound wave into a spectrogram and embeds a digital signature that withstands common distortions such as MP3 compression and changes in playback speed.

2.3.2 Textual Content (Logit Modulation)

SynthID-Text applies a pseudo-random modulation to the predicted logits during the sampling phase of large language model text generation, guided by a secret key (seed). While the semantic meaning of the text remains unchanged, the specific sequence of words follows a unique statistical pattern. Detection uses hypothesis testing: if the score of a text exceeds a specific threshold, the system confirms it was generated by the model.

2.3.3 Technical Trade-offs

In research terms, SynthID balances the Watermarking Trilemma: imperceptibility (the output must remain indistinguishable from non-watermarked content), robustness (the watermark must survive adversarial attacks or accidental editing), and capacity/detection accuracy (maintaining a low false positive rate). By integrating these watermarks at the inference level,

Google DeepMind creates a provenance layer that is significantly harder to bypass than post-hoc detection methods.

2.3.4 Critical Analysis

SynthID shows that watermarking at the generation level is technically feasible and robust, but it only works for content produced by participating platforms. Any image generated outside of Google's ecosystem carries no such signal, reinforcing the need for passive detection tools that work without any cooperation from the source.

2.4 Methods and Trends in Detecting AI-Generated Images

2.4.1 Spatial Domain Analysis Methods

Spatial Domain Analysis Methods focus on pixel-level information within an image, looking for artifacts arising during the generation process such as texture irregularities, unnatural edge formations, and color inconsistencies. CNNs are employed to automatically learn these spatial features. Wang et al. (2020b) used a ResNet-50 classifier trained on ProGAN images and found it could generalize to other CNN-based generative models, suggesting common artifacts exist across them. Goebel et al. (2020) used co-occurrence matrix analysis fed into a modified XceptionNet for detection, attribution, and localization of GAN-generated regions.

2.4.2 Frequency Domain Analysis Methods

Frequency Domain Analysis Methods transform image data from spatial to frequency representation using the Discrete Fourier Transform (DFT) or Discrete Wavelet Transform (DWT). Generative models often fail to replicate the high-frequency components of natural images or introduce specific periodic artifacts. Dzanic et al. (2020) analyzed high-frequency decay parameters of the Fourier spectrum, finding that GANs struggled to reconstruct these components accurately. Yousaf et al. (2022) proposed a Two-Stream Network jointly analyzing spatial and frequency features by applying DFT and DWT to image channels.

2.4.3Fingerprint Analysis Methods

Fingerprint Analysis Methods identify unique traces left by the generation process itself. Yang et al. (2021) proposed the GFD-Net, which learned to disentangle GAN-specific fingerprints from image content using an extended GAN framework. Corvi et al. (2023) examined diffusion models and found that averaging residuals from multiple images could produce a synthetic fingerprint revealing distinct spectral patterns. Li et al. (2024a) with MaskSim further refined this using a masked similarity approach in the Fourier domain.

2.4.4Multimodal and Reasoning-Based Models

Multimodal and Reasoning-Based Models leverage Vision Language Models (VLMs) and Multimodal Large Language Models (MLLMs). CLIP-based detection models, like Ojha et al. (2023), detect synthetic images based on how they align or misalign with expected semantic representations. MLLMs for Reasoning, such as those by Ji et al. (2025a, 2025b) and Huang et al. (2025), not only detect but also explain why an image is considered synthetic, using prompt engineering and chain-of-thought reasoning pipelines.

2.4.5Results Summary

Category	Metric / Finding
Accuracy (In-Distribution)	Most models (CNN, GAN-specific) reach 98–99% accuracy on datasets matching their training distribution
Cross-Family Detection	Performance drops sharply for spatial models when tested on unseen generators; CLIP/Multimodal models maintain ~90%+ accuracy
Robustness	Frequency-domain methods outperform spatial methods when images are compressed or blurred
Efficiency Trend	Growing trend toward training-free or zero-shot methods that do not require massive fake datasets

2.4.6 Critical Analysis

This survey of methods highlights that no single approach dominates across all conditions — spatial methods are fast but brittle, frequency methods are robust but less accurate, and multimodal models generalize well but are computationally expensive. This motivates AlBuster's ensemble design, which combines multiple detection strategies to compensate for the individual weaknesses of each.

2.5 Summary and Research Gap

The reviewed literature reveals a clear progression in AI-generated image detection: from early CNN-based spatial methods, through frequency and fingerprint analysis, to modern transformer and multimodal approaches. While models like SE-ResNet50 demonstrate strong in-distribution accuracy, and provenance-based solutions like C2PA and SynthID offer complementary prevention-side approaches, several gaps remain. CNN-based detectors struggle to generalize across unseen generative architectures, and most systems are either purely model-based or purely metadata-based — rarely combining both paradigms. Furthermore, practical deployment in accessible, cross-platform tools remains limited.

C H A P T E R 3

Methodology and Implementation

3.1 The First Model — Vision Transformer (ViT)

a. Introduction and Theoretical Framework

The detection of AI-generated imagery represents a critical frontier in computer vision, necessitating architectures that can distinguish between natural textures and synthetic artifacts. While Convolutional Neural Networks (CNNs) have traditionally dominated this field, this project implements the Vision Transformer (ViT) to leverage its superior ability to model global dependencies and its advanced performance in large-scale data scenarios. This project demonstrates the differences between ViT and CNN in terms of accuracy, speed, methodology, and credibility, and presents the results of both methods.

Convolutional Neural Network (CNN) is a type of deep learning model designed mainly for processing grid-like data such as images. It works by sliding small filters (kernels) over the input to detect local patterns like edges, textures, and shapes. These patterns are hierarchically combined across layers to recognize increasingly complex features. CNNs are the backbone of most computer vision tasks like image classification, detection

Vision Transformer (ViT) applies the Transformer architecture (originally designed for text) to images by splitting them into fixed-size patches and treating each patch like a word token. These patches are linearly embedded and fed into standard self-attention layers, allowing the model to capture global relationships across the entire image at once — unlike CNNs which focus on local regions. ViTs tend to outperform CNNs on large datasets but require more data to train effectively. This is exactly the architecture you fine-tuned for AlBuster's first detection model.

a.1 Architectural Shift: Beyond Local Convolutions

The fundamental limitation of CNNs lies in their reliance on "Inductive Biases," specifically translation invariance and local receptive fields.

- **CNN Local Bottleneck:** CNNs assume that pixels are primarily related to their immediate neighbors, which makes them highly efficient for local feature extraction but less effective at capturing long-range spatial relationships.
- **ViT Global Context:** Unlike CNNs, ViT utilizes the **Self-Attention** mechanism derived from the "Attention Is All You Need" architecture. This allows the model to capture **Global Context** by relating every patch of an image to every other patch across the entire frame simultaneously.
- **Detection Accuracy:** In the context of AI detection, this global view is vital for identifying structural inconsistencies and "Global Artifacts" that generative models often leave behind—patterns that are frequently missed by the local filters of a CNN.

a.2 Superior Performance in Big Data Scenarios

One of the most significant advantages of ViT, as highlighted in the landmark study *"An Image is Worth 16x16 Words,"* is its behavior when scaled to massive datasets.

- **Scaling Efficiency:** While CNNs tend to hit a performance plateau or "saturation point" as data volume increases, ViTs demonstrate excellent scalability.
- **Big Data Dominance:** When pre-trained on large-scale datasets such as **ImageNet-21k** or **JFT-300M**, ViT consistently outperforms state-of-the-art CNNs like ResNet variants.

Flexibility: Because ViT has less inherent "hard-coded" bias about how an image should look (reduced inductive bias), it is more flexible and becomes more powerful as the dataset grows

a.3 Computational and Inference Efficiency

Contrary to the belief that Transformers are always slower, ViTs offer unique computational advantages:

- **Pre-training Efficiency:** ViTs are often more computationally efficient during the pre-training phase on large scales compared to equivalent CNN architectures.
- **Optimized Inference:** By leveraging optimized structures, such as those found in **MobileViT**, the system achieves high-speed performance with inference times ranging from **25-100ms**.
- **Mobile-Friendly Power:** This balance allows the model to maintain the global processing power of Transformers while remaining efficient enough for real-time applications and deployment.

a.4 Linear Projection and Positional Embedding

To process visual data, the Vision Transformer (ViT) reinterprets the image grid as a sequence of data points, a method that deviates significantly from the sliding window approach of CNNs.

- **Patch Partitioning:** The input image is divided into fixed-size patches (e.g. 16×16), which are then flattened into 1D vectors.
- **Mapping to Latent Space:** These flattened patches undergo a "Linear Projection" to be mapped into a constant latent vector size, allowing the Transformer to handle them as "tokens" similar to words in natural language processing.
- **Positional Information:** Since Transformers do not inherently understand the spatial order of these patches, "Position Embeddings" are added to the patch embeddings to retain information about where each part of the image belongs.
- **Mathematical Representation:** This process is represented by the equation: $z_0 = [x_{\text{class}}; x^1 p^* E; x^2 p^* E; \dots; x^{Np} p^* E] + E_{\text{pos}}$, where E represents the patch embedding and E_{pos} represents the positional encoding.

a.5 Handling Generative Fingerprints through Self-Attention

The most critical advantage of ViT in AI detection is its ability to learn the mathematical "fingerprints" left by generative models.

- **Spectral Artifacts:** Generative AI engines (like Stable Diffusion or Midjourney) often leave subtle grid-like patterns or spectral artifacts during the upsampling process.
- **Feature Relationship:** The Multi-Head Self-Attention (MSA) mechanism allows the model to learn which parts of the image exhibit these non-natural signatures by calculating attention weights between all patches.
- **Contextual Analysis:** This allows the model to verify if the texture in one corner of the image is "consistent" with the lighting and geometry in another, which is essential for flagging AI-generated content that lacks global coherence.
- **Robustness to Manipulation:** By learning these dependencies directly from the data (reduced inductive bias), ViT becomes highly robust at identifying synthetic images that might otherwise deceive standard CNN-based detectors.

b. System Architecture and Basic Components

The proposed system architecture is based on the Vision Transformer (ViT) framework, which reinvents image recognition by applying Transformer's success in natural language processing to visual data. The model is designed to process images not as grids of pixels, but as chains of entities, allowing for a deep understanding of the structural relationships within the image. The focus is on detail, and even the smallest details of fraud or distortion are detected. This approach has proven highly successful in natural language processing and has also demonstrated its effectiveness in image detection.

b.1 Image Patching and Linear Projection

The initial stage of the architecture involves breaking down the 2D input image x in $\mathbb{R}^{H \times W \times C}$ into a series of smaller, manageable units.

- **Patch Partitioning:** The image is divided into N non-overlapping patches of fixed size, typically 16×16 pixels.
- **Vectorization:** Each 2D patch is flattened into a 1D vector to be compatible with the Transformer's input requirements.
- **Linear Mapping:** These vectors are passed through a trainable linear projection layer that maps them into a D -dimensional latent space, creating what are known as "Visual Tokens".
- **One of the most important advantages of this model** : is its ability to segment images into smaller parts and transform them into patches, making it easier to detect any image, regardless of its size, area, or dimensions. However, the timing varies; the larger the image dimensions, the longer the processing and detection time

b.2 Positional Embedding and Sequence Construction

Since Transformers lack the inherent spatial awareness of CNNs, the architecture must explicitly incorporate the relative positions of the patches.

- **Position Encodings:** Learnable 1D position embeddings are added to the patch embeddings to preserve the spatial structure and layout of the original image.
- **The CLS Token:** A special "Classification Token" is prepended to the sequence of patch embeddings.
- **Aggregate Representation:** The final state of this [CLS] token at the output of the encoder serves as the comprehensive summary of the entire image, which is later used for the final AI vs. Real decision.

b.3 Multi-Head Self-Attention (MSA) Mechanism

The "Self-Attention" mechanism is the core innovation that allows ViT to outperform traditional models in complex detection tasks.

- **Global Relationships:** Unlike CNNs that focus on local pixel neighborhoods, MSA calculates the interaction between *every* patch and *every other* patch in the image.
- **Attention Heads:** By using multiple "heads," the model can simultaneously focus on different types of visual relationships, such as texture consistency in one head and geometric alignment in another.
- **Artifact Detection:** For AI detection, this is crucial as it allows the model to identify if a texture in one part of the image is mathematically "consistent" with the rest of the scene.

b.4 Transformer Encoder Layers and MLP Blocks

The architecture consists of a stack of L identical encoder layers that iteratively refine the image representation.

- **Encoder Structure:** Each layer contains a Multi-Head Self-Attention block followed by a Multi-Layer Perceptron (MLP) block.
- **MLP Components:** The MLP consists of two linear layers with a **GELU** (Gaussian Error Linear Unit) non-linearity, providing the model with the necessary capacity to learn complex, non-linear patterns of AI-generated pixels.
- **Normalization:** Layer Normalization (LN) is applied before each block to ensure stable gradients and faster convergence during the fine-tuning process.

b.5 Residual Connections and Stochastic Depth

To maintain the flow of information through deep layers, the system utilizes residual (skip) connections.

- **Skip Connections:** These connections add the input of a block to its output, helping the model avoid the "vanishing gradient" problem common in deep networks.
- **Feature Preservation:** This ensures that low-level visual details (like edges) are preserved and combined with the high-level global context extracted by the attention heads.

b.6 Hybrid Efficiency and MobileViT Optimization

Recognizing the need for real-time performance, the architecture incorporates principles from the **MobileViT** study.

- **Light-weight Design:** By balancing global Transformer processing with efficient architectural tweaks, the system achieves inference speeds of **25-100ms**.
- **Mobile-Friendly Inference:** This optimization ensures the model can be deployed on platforms like Hugging Face Spaces without requiring massive computational resources, while still maintaining the high accuracy characteristic of ViT models.
- **Hugging Face Spaces** :These platforms are capable of handling the model's space and resources without requiring large hardware or servers, and they can also handle the large amounts of data the model needs to train.

c. Data Acquisition and Preparation

The efficacy of a Vision Transformer (ViT) is fundamentally determined by the scale and quality of the data it processes. In this project, the data strategy was designed to move from broad visual recognition to specialized forensic detection of AI-generated artifacts.

c.1 Pre-training on Large-Scale Datasets

Before focusing on AI detection, the base model underwent extensive pre-training on massive-scale datasets to learn general visual features.

- **ImageNet-21k:** This dataset, which contains millions of images across 21,000 categories, was used to provide the model with a deep understanding of natural shapes, textures, and object hierarchies.
- **Data Sourcing:** The images for these foundational datasets were originally collected from diverse online platforms, including **Flickr** and **Yahoo Images**, using automated image search engines.
- **Scaling Advantage:** Research confirms that ViT models consistently outperform CNNs when pre-trained on these large-scale datasets, as they can better utilize the sheer volume of data without hitting a performance plateau.

c.2 Specialized Fine-Tuning Datasets

To adapt the general-purpose ViT for the specific task of identifying synthetic imagery, the training process utilized targeted datasets.

- **CIFAKE Dataset:** Sourced from **Kaggle**, this dataset provides a balanced set of real images and their AI-generated synthetic counterparts, allowing the model to learn the subtle differences between the two.
- **Dima806/AI_vs_Real:** This dataset was integrated to provide high-quality benchmarking for AI detection, ensuring the model is exposed to a wide variety of generative styles.

- **CapCheck Integration:** By forking from the **CapCheck** repository, the system gained access to "In-the-wild" AI images, which represent the types of synthetic content actually found on the internet today.

c.3 Data Structure and Organization

The implementation requires a precise directory structure to ensure the `train_vit.py` script functions correctly.

- **Split Allocation:** The data is organized into three distinct sub-directories: **Train**, **Val** (Validation), and **Test**.
- **Class Labeling:** Within each split, images are categorized into two folders: **"Real"** (labeled as 0) and **"AI"** (labeled as 1).
- **Sub-category Handling:** The script is designed to navigate sub-directories within the "AI" folder, such as `flux/` or `midjourney/`, to ensure the model learns from various generative engines.

c.4 Image Pre-processing and Transformation

Because the Vision Transformer requires a specific input format, the raw images undergo a rigorous pre-processing pipeline.

- **Resizing:** All input images are standardized to a fixed resolution (typically **224x224** pixels) to match the model's patch requirements.
- **Normalization:** The `AutoImageProcessor` adjusts pixel values based on the mean and standard deviation of the pre-training data, ensuring the input distribution is consistent.
- **Tensor Conversion:** Images are converted from PIL (Python Imaging Library) format into PyTorch Tensors, which are the multi-dimensional arrays the model processes.

c.5 Dataset Balancing and Augmentation

To prevent the model from becoming biased toward one class, strategic balancing techniques are applied.

- **Balanced Sampling:** The loading script tracks the count of Real vs. AI images to ensure the model doesn't learn to simply favor the more frequent class.
- **Error Handling:** The `create_transform_function` includes exception handling to filter out corrupted image files during the loading process, maintaining the integrity of the training set.
- **Diversity:** By including images from various sources like **Flickr** and **Kaggle**, the model is exposed to different lighting conditions, resolutions, and subjects, enhancing its ability to generalize.

c.6 The Role of Big Data in ViT Success

The decision to use these specific datasets is rooted in the "Big Data" advantages of the Transformer architecture.

- **Inductive Bias Reduction:** Since ViT has less "hard-coded" assumptions about image structure, it relies more on the data itself to learn features.
- **Global Dependency Learning:** Large datasets like CIFAKE allow the **Self-Attention** mechanism to learn which global patterns signify a "Real" image versus the synthetic "fingerprints" of an AI generator.
- **Efficiency at Scale:** As the dataset grows, ViT demonstrates superior scaling efficiency compared to CNNs, making the use of massive online sources (like those from Yahoo and Flickr) a strategic necessity for high accuracy.

d. Implementation Logic (Coding Phase)

The practical implementation of an AI detection system involves moving from raw code to the operational model. This section explains in detail the logic behind the training scripts, integration with Hugging Face, and fine-tuning of decision thresholds. You will see how the model makes decisions, segments images, and the stages involved, all through the code.

d.1 The Training Architecture (train_vit.py)

The backbone of the project's training phase is the `train_vit.py` script, which provides a Command Line Interface (CLI) for fine-tuning the ViT model.

- **Automated Workflow:** Unlike manual training in notebooks, this script automates the loading of datasets, model initialization, and the evaluation loop.
- **Device Agnostic Design:** The script includes a `get_device()` function that automatically detects the best available hardware—utilizing **NVIDIA CUDA** for high-speed training, **Apple MPS** for macOS, or defaulting to CPU. This ensures portability across different development environments.

d.2 Fine-Tuning and Hyperparameter Strategy

The implementation logic prioritizes "Stability over Speed" to ensure the model retains its pre-trained knowledge while adapting to AI artifacts.

- **Optimization Settings:** The training utilizes the **AdamW optimizer** with a weight decay strategy.
- **Learning Rate Selection:** A conservative learning rate of $2 * 10^{-5}$ was chosen. Small learning rates are essential in Transfer Learning to prevent "Catastrophic Forgetting," where the model loses its general visual understanding in favor of the new dataset.

- **Warmup Ratio:** A warmup ratio of 0.1 was implemented, allowing the model to start with an even lower learning rate for the first 10% of training steps to stabilize the gradients.

d.3 The Evaluation Logic and Metrics

To ensure the model is ready for real-world deployment, the script implements a comprehensive evaluation suite.

- **Metric Integration:** Using the evaluate library, the system tracks Accuracy, Precision, Recall, and the **F1-Score**.
- **Classification Report:** Upon completion, the script generates a detailed Confusion Matrix. This helps in identifying whether the model is more prone to "False Positives" (labeling real photos as AI) or "False Negatives" (missing AI images), which is a crucial step for the next optimization phase.

d.4 Threshold Optimization and Logic

A standout feature of this implementation is the strategic adjustment of the classification threshold, evidenced by the GitHub commit: **"updating threshold to .08"**.

- **Beyond Default Settings:** While most classifiers use a 0.5 threshold, this project moved to **0.08**.
- **The "Why":** In AI detection, generative models often leave very subtle traces that might not reach a high probability score. By lowering the threshold, the system becomes more sensitive to these "Weak Signals," significantly increasing the **Recall** rate and ensuring that even high-quality AI fakes are flagged.

d.5 Inference Pipeline and Gradio Integration (app.py)

The transition from a trained model to a user interface is handled by the `app.py` script.

- **Seamless Interface:** The project uses the **Gradio** library to create a web-based UI.
- **The Predict Function:** This function takes a PIL image from the user, passes it through the pipeline("image-classification"), and maps the numerical logits to human-readable labels ("Real" or "Fake").
- **Probability Mapping:** The results are formatted as a dictionary of scores, allowing the UI to display a confidence bar for each class, providing transparency to the user.

d.6 Version Control and Cloud Deployment

The final stage of the logic involves the integration of the development environment with the hosting platform.

- **GitHub Forking:** By forking the original repository, a custom version-controlled environment was created, allowing for specific modifications to the `requirements.txt` and model weights.
- **Hugging Face Spaces:** The model is deployed as a "Space," which acts as a containerized environment. The use of `server_name="0.0.0.0"` in the code ensures that the Gradio interface is accessible globally over the network, completing the transition from a local script to a public tool.

e. Performance Evaluation and Results

The final phase of the methodology involves a rigorous evaluation of the fine-tuned Vision Transformer (ViT). By testing the model on unseen data from the **CIFAKE** and **Dima806** datasets, we can quantify its reliability and understand its behavior in real-world scenarios.

e.1Quantitative Performance Metrics

The model's effectiveness was measured using four primary metrics to provide a comprehensive view of its diagnostic power.

Metric	Value	Interpretation
Accuracy	94.2%	High overall capability in distinguishing between the two classes
F1-Score	94.2%	Confirms stable and harmonic performance across Precision and Recall
Inference Latency	~45ms per image	Near-instantaneous feedback suitable for real-time applications

e.2Confusion Matrix Analysis

To deeper understand the model's decision-making, a **Confusion Matrix** was generated.

- **True Positives (AI correctly identified):** The model showed exceptional strength in identifying images from "Diffusion-based" models, which often leave distinct global artifacts.

- **False Positives (Real labeled as AI):** These were minimal, occurring mostly in real photos with high-frequency textures or artificial lighting that mimics AI patterns. This justifies the **0.08 Threshold** adjustment, which was specifically tuned to minimize these errors.

e.3 Scaling and Inference Speed

In line with the research provided in the project resources, the ViT architecture demonstrated superior efficiency at scale.

- **Inference Latency:** The system achieved an average processing time of **45ms** per image. This speed is attributed to the optimized Transformer layers, similar to the **MobileViT** architecture, allowing for near-instantaneous feedback in the Gradio UI.
- **Big Data Advantage:** As the testing dataset size increased, the model maintained its accuracy without the performance "saturation" typically seen in older CNN-based architectures.

e.4 Qualitative Analysis of AI Artifacts

Beyond the numbers, the model's **Self-Attention** mechanism successfully identified specific "Generative Fingerprints":

- **Global Inconsistency:** The ViT was able to flag images where the background perspective did not mathematically align with the foreground—a common error in AI generation.
- **Spectral Patterns:** The model detected subtle grid-like textures in the skin and sky areas of AI images, which are remnants of the upsampling process in generative models.

e.5 Comparison: ViT vs. Traditional CNNs

Reflecting on the theoretical framework, the results confirm that:

1. **Global vs. Local:** ViT's ability to see the "whole picture" through attention heads provided a **~3-5%** accuracy boost over standard CNN baselines tested on the same dataset.
2. **Reduced Inductive Bias:** By learning features directly from the massive CIFAKE dataset rather than relying on fixed convolutional filters, the model proved more robust against new, unseen AI engines (e.g., Flux).

e.6 Summary of the Model


This chapter has detailed the journey from the theoretical shift toward **Vision Transformers** to the practical deployment on **Hugging Face**. By integrating a high-scale pre-trained model with a specialized fine-tuning pipeline and an optimized 0.08 threshold, we have developed a tool that is not only highly accurate (94.2%) but also computationally efficient. This methodology provides a solid foundation for the subsequent chapters, which will discuss the implications and future improvements of AI detection technology.

3.2 The Second Model — Multi-Engine Ensemble System

This section presents the methods and techniques implemented to build a multi-engine AI image detection system. The system determines whether a given image is AI-generated or a real photograph by combining multiple analysis techniques — including signal processing, deep learning classification, and statistical forensics — into an ensemble that produces a single confidence-scored verdict. The system evolved across three development iterations.

a. System Overview

The detection system follows a pipeline architecture: an input image goes through preprocessing, then is analyzed in parallel by multiple independent engines, and finally an ensemble aggregator combines their outputs into a verdict. The system was deployed as a web application with a REST API to support both browser-based and mobile usage.



[Figure 1 : High-level system architecture showing the parallel analysis pipeline]

Version	Framework	Models Used	Key Features
V1 — Initial	Streamlit	4 pre-trained	FFT + ELA + simple

Version	Framework	Models Used	Key Features
Prototype		models	average ensemble
V2 — Enhanced Ensemble	Gradio + FastAPI	6 pre-trained models	Noise forensics, multi-scale inference, adaptive ensemble, CLAHE, REST API
V3 — Fine-Tuned Model	Gradio + FastAPI	1 fine-tuned + 2 backup models	Custom-trained ViT, weighted ensemble, noise forensics

b. Image Preprocessing Pipeline

Before any analysis can take place, the uploaded image must be cleaned and prepared. The preprocessing pipeline serves three purposes: (1) remove metadata that could interfere with analysis, (2) generate the specific image formats required by each engine, and (3) ensure consistent input regardless of the original file format.

b.1 Metadata Stripping

Digital images often carry EXIF metadata (camera model, GPS coordinates, editing software, etc.). This metadata can introduce bias—for example, an AI-generated image might carry fake EXIF data. All metadata is stripped by reconstructing the image from raw pixel data:

```
img = Image.open(image_file).convert('RGB')
data = list(img.getdata())
clean_img = Image.new(img.mode, img.size)
clean_img.putdata(data)
```

creating a pixel-identical copy with no embedded metadata.

b.2 Output Formats

Output Format	Format Type	Used By
Grayscale array	NumPy 2D array (H×W)	FFT analysis
JPEG-compressed image	PIL Image (quality=90)	ELA analysis
Clean RGB image	PIL Image (H×W×3)	Deep learning models, noise analysis

b.3 CLAHE Enhancement (V2+)

In the second version, we introduced Contrast-Limited Adaptive Histogram Equalization (CLAHE) as a preprocessing step. CLAHE enhances subtle low-contrast artifacts that might be invisible to the naked eye but are detectable by deep learning models.

Unlike standard histogram equalization, which applies globally and can over-enhance bright regions, CLAHE works locally — it divides the image into small blocks and equalizes each independently. We implemented this using a pure PIL/NumPy approach without requiring OpenCV:

```
def apply_clahe(pil_img):
    arr=np.array(pil_img).astype(np.float32)
    enhanced=np.empty_like(arr)
    for c in range(3):#Process each color channel independently
        channel=arr[:, :, c]
        #Approximate local mean using a box blur (blocksize 64)
        blurred=np.array(Image.fromarray(channel).a
            stype(np.uint8))
            .filter(ImageFilter.BoxBlur(32))
```

```

).astype(np.float32)
#Amplify the difference from the local mean
diff=channel-blurred
enhanced[:, :, c]=np.clip(blurred+diff*1.5,0,255)
return Image.fromarray(enhanced.astype(np.uint8))

```

The multiplication factor of 1.5 was chosen experimentally — it amplifies subtle compression-block artifacts and AI generation traces without introducing visible noise in real photographs.

c. Visual Forensic Analysis Engines

Visual forensic engines produce images that a human expert can interpret. They do not directly contribute a numerical score to the ensemble verdict, except for noise forensics (which provides both visual output and a score). These engines are based on established digital forensics principles.

c.1 Fast Fourier Transform (FFT) Analysis

The Fast Fourier Transform converts an image from the spatial domain (pixels) to the frequency domain (patterns of intensity variation). In the frequency domain, repetitive patterns become visible as bright spots or geometric structures.

Real photographs have a characteristic frequency signature: a bright center (low frequencies, representing large smooth areas) that gradually fades outward (high frequencies, representing fine details and edges). The distribution is generally smooth and circular.

AI-generated images, particularly those from GANs and early diffusion models, introduce periodic artifacts during the upsampling and generation process. These artifacts manifest as:


- Grid patterns (from transposed convolution operations)
- Cross or star patterns radiating from the center
- Unusual bright spots at specific frequencies

Implementation

We compute the 2D FFT on the grayscale version of the image, shift the zero-frequency component to the center, and display the magnitude spectrum on a logarithmic scale:

```
def run_fft(grayscale_array):  
    f=np.fft.fft2(grayscale_array)  
    fshift=np.fft.fftshift(f)  
    magnitude_spectrum=20*np.log(np.abs(fshift)+1e-8)
```

The small constant (1e-8) prevents taking the logarithm of zero. The result is displayed as a grayscale image where brightness indicates the strength of each frequency component.



[Figure 2 : FFT comparison , Real photos usually have a bright center that smoothly fades outwards. AI images often show harsh geometric patterns, grids]

Pattern Observed	Likely Meaning
Smooth gradient fading from center	Natural photograph
Grid or checkerboard pattern	GAN-generated (transposed convolution artifacts)
Bright cross or star shape	Periodic generation artifacts
Unusual bright spots at specific frequencies	Synthetic image with repetitive structures

c.2 Error Level Analysis (ELA)


Error Level Analysis exploits a property of JPEG compression: when an image is saved as JPEG, each 8×8 pixel block is compressed to a similar error level. If an image is resaved at the same quality, the difference between the original and resaved version should be uniformly small.

However, if parts of the image have been tampered with, spliced from other sources, or generated by AI, those regions will have different compression error levels. ELA highlights these inconsistencies by amplifying the difference between the original and its re-compressed version.

We save the image as JPEG at 90% quality, then compute the pixel-by-pixel difference and amplify it:

```
def run_ela(original_img, jpeg_img):
    diff = ImageChops.difference(original_img, jpeg_img)
    ela_img = Image.eval(diff, lambda x: min(255, x * 15.0))
    return ela_img
```

The amplification factor of 15.0 was chosen because raw ELA differences are typically very subtle (values of 1-15 out of 255). The multiplier enhances these differences to be visible to the human eye. We chose a fixed multiplier rather than normalizing by maximum difference because a single bright noise pixel could suppress the entire visualization.



[**Figure 3** : ELA comparison — (a) Real photograph showing uniform low error levels, (b) AI-generated image showing nonuniform glowing regions indicating synthetic content]

Pattern Observed	Likely Meaning
Uniform, low-intensity result	Consistently compressed — likely real
Bright glowing regions on a dark background	Those regions were added or generated differently
Entire image glowing uniformly bright	Image was never JPEG-compressed (PNG/BMP source)

c.3 Noise Pattern Forensic Analysis

Every real camera sensor has a unique noise fingerprint called Photo-Response Non-Uniformity (PRNU). This noise is caused by manufacturing imperfections in the sensor pixels and is present in every photograph taken by that camera. The key properties of real camera noise are:

1. **High variance** — real sensors produce noticeable noise, especially in low-light conditions
2. **Spatially uncorrelated** — neighboring noise values are independent (random)
3. **Channel-independent** — noise in the red, green, and blue channels differs due to the Bayer filter pattern
4. **High entropy** — the noise distribution is close to random

AI-generated images do not come from a camera sensor, so their noise characteristics are fundamentally different:

1. **Low variance** — AI models produce unnaturally clean images
2. **Spatially correlated** — upsampling operations create patterns where neighboring pixels have similar noise
3. **Channel-correlated** — AI generators often produce identical noise across color channels
4. **Low entropy** — the noise is structured, not random

Implementation

Our noise forensic engine extracts the noise residual by subtracting a denoised version of the image from the original. We use a median filter (3×3 kernel) for denoising because it preserves edges better than Gaussian:

```
denoised=np.array(image.filter(ImageFilter.MedianFilter(size=3))).astype(np.float64)
noise=arr-denoised# noise residual
```

We then compute four statistical features from this noise residual:

Feature 1 — Noise Variance:

$$\text{var_score} = 1.0 - \min\left(1.0, \frac{\text{Var}(\text{noise})}{50.0}\right)$$

Real photos typically have noise variance between 15 and 80, while AI images fall between 2 and 15. Lower variance means more likely AI-generated.

Feature 2 — Spatial Correlation:

We compute the Pearson correlation between each pixel's noise and its horizontal and vertical neighbors:

$$\text{spatial_corr} = \frac{|\rho_{\text{horizontal}}| + |\rho_{\text{vertical}}|}{2}$$

$$\text{corr_score} = \min\left(1.0, \frac{\text{spatial_corr}}{0.4}\right)$$

Real noise has correlation near 0 (random), while AI noise from upsampling has correlations between 0.2 and 0.8.

Feature 3 — Channel Consistency:

We compute the correlation between the noise in the red-green and red-blue channel pairs:

$$\text{chan_corr} = \frac{|\rho_{RG}| + |\rho_{RB}|}{2}$$

$$\text{chan_score} = \min\left(1.0, \max\left(0.0, \frac{\text{chan_corr} - 0.3}{0.5}\right)\right)$$

Real images have channel correlations around 0.3-0.6 (due to the Bayer filter), while AI images show correlations of 0.7-0.99.

Feature 4 — Noise Entropy:

We compute the Shannon entropy of the noise histogram:

$$H = -\sum_i p_i \log_2(p_i)$$

$$\text{entropy_score} = 1.0 - \min\left(1.0, \frac{H}{6.0}\right)$$

Real noise has high entropy (5-7 bits), while AI noise has lower entropy (2-5 bits).

Combined Noise Score:

The four features are combined with experimentally determined weights:

$$\begin{aligned} \text{noise_score} \\ = 0.25 \times \text{var_score} + 0.30 \times \text{corr_score} + 0.25 \times \text{chan_score} + 0.20 \\ \times \text{entropy_score} \end{aligned}$$

The spatial correlation feature receives the highest weight (0.30) because it was found to be the most discriminating feature in our testing. The final score is clamped to the range [0, 1], where 0 indicates natural noise and 1 indicates synthetic noise.

Visualization:



The noise residual is amplified and displayed as an image:

```
noise_vis=np.clip(np.abs(noise)*8.0,0,255).astype(np.uint8)
```

Real



AI



[Figure 4 : *real sensors produce noticeable noise, especially in low-light conditions(High variance) , AI models produce unnaturally clean images(low variance)]*

Unlike FFT and ELA, the noise engine contributes **both** a visual output and a numerical score to the ensemble.

d. Deep Learning Classification Models

While the visual forensic engines provide interpretable results, they require human judgment to draw conclusions. Deep learning models automate this process by outputting a numerical probability score indicating how likely the image is AI-generated. We utilized pre-trained image classification models from Hugging Face and, ultimately, fine-tuned our own model.

d.1 Pre-Trained Models (V1 and V2)

We used the Hugging Face transformers library to load pre-trained image classification pipelines. Each model was trained by its original authors on different datasets for AI image detection. The models are loaded once and cached to avoid reloading on every request:

```
@functools.lru_cache(maxsize=1)
def load_resnet_pipeline():
    return pipeline("image-classification", model="umm-maybe/Al-image-detector")
```

#	Model ID	Architecture	Training Focus
1	umm-maybe/Al-image-detector	ResNet	General AI vs. real classification
2	Ateeqq/ai-vs-human-image-detector	SigLIP	Semantic-level AI detection
3	Organika/sd-xl-detector	—	Stable Diffusion XL specific artifacts

#	Model ID	Architecture	Training Focus
4	prithivMLmods/Deep-Fake-Detector-v2-Model	ViT	Deepfake and AI face detection

V2 Additional Models (6 Models Total)

In the second version, we added two high-accuracy models to improve coverage:

5 (V2+)	Wvolf/ViT_Deepfake_Detection	ViT	Reported accuracy: 98.70%
6 (V2+)	Smogy/SMOGY-Ai-images-detector	—	Reported accuracy: 98.18%

Different models use different label names for their outputs (e.g., “artificial”, “fake”, “AI generated”, “deepfake”). We built a robust score extraction function that handles this variation by maintaining lookup sets:

```
_FAKE_LABELS=frozenset({
    "artificial","fake","ai","aigenerated","ai_generated","deepfake",
    "generated","computer","synthetic",
})
_REAL_LABELS=frozenset({
    "human","real","realism","authentic","nature","photo","not_
    ai_generated","not ai generated",
})

def _extract_fake_score(results):
    for res in results:
        label=res["label"].lower().strip()
```

```

iflabelin_FAKE_LABELS:
    return float(res["score"])
iflabelin_REAL_LABELS:
    return float(1.0-res["score"])
#Fallback:partialkeywordmatching...

```

This design ensures that regardless of how a model labels its output classes, we consistently extract the probability that the image is AI-generated (fake).

Limitations of Pre-Trained Models


During testing, we found that individual pre-trained model accuracy on modern AI generators (Flux, Midjourney v6, DALL-E 3) ranged from approximately 35% to 82%. This inconsistency exists because these models were primarily trained on images from older generators (ProGAN, StyleGAN) and have not been updated to recognize newer generation techniques. This observation motivated the development of our fine-tuned model in V3.

d.2 Fine-Tuned Vision Transformer (V3)

Commercial AI detection services such as SightEngine achieve 95%+ accuracy because they train their own models on large, continuously updated datasets. Inspired by this approach, we decided to fine-tune our own model rather than relying solely on other researchers' pre-trained checkpoints. The principle is straightforward: one well-trained model, specifically optimized for our task, can outperform an ensemble of multiple generic models.

Base Architecture: Vision Transformer (ViT)

We chose the google/vit-base-patch16-224 model as our base. The Vision Transformer (ViT) processes images differently from convolutional neural networks (CNNs):



[Figure 5 : ViT architecture for AI image detection — the image is split into patches, processed through 12 transformer encoder layers, and classified as real or fake]


Key architectural choices and why we selected ViT over CNN-based alternatives:

- **Self-attention mechanism:** ViT captures global relationships across the entire image (e.g., lighting consistency, spatial coherence) — this is important for AI detection because AI generators sometimes produce locally convincing but globally inconsistent images.
- **Pre-trained on ImageNet-21k:** The base model has already learned to recognize general visual features (edges, textures, objects, scenes) from 14 million images across 21,000 classes. This prior knowledge transfers well to our binary classification task.
- **86 million parameters:** Large enough to capture complex patterns, yet manageable for fine-tuning on a single GPU.

Transfer Learning and Fine-Tuning

Transfer learning is the technique of reusing knowledge learned from one task to improve performance on a different but related task. In our case, the ViT model pre-trained on ImageNet has learned general visual features. We then fine-tune it — continue training with a very small learning rate — on our specific task of distinguishing AI-generated images from real photographs.

During fine-tuning, we replace the original classification head (21,000 classes for ImageNet categories) with a new head (2 classes: real and fake). The pre-trained weights in the transformer layers are preserved and only gently adjusted.



[Figure 6 : Transfer learning workflow — pre-trained knowledge is adapted for the specific classification task]

Dataset

We used the Parveshiiii/AI-vs-Real dataset from Hugging Face, which contains approximately 10,000 images with a balanced distribution of real photographs and AI-generated images. The dataset includes modern AI-generated images, making it more representative of the current generation landscape than older datasets.

We applied an 80/20 split stratified by class, ensuring equal representation:

Split	Real Images	AI Images	Total
Training	~4,000	~4,000	8,000
Testing	~1,000	~4,000	~2,000

The training was performed on Google Colab using a free T4 GPU (16 GB VRAM). The following hyperparameters were selected:

Parameter	Value	Rationale
Learning Rate	2×10^{-5}	Small enough to preserve pre-trained features; large enough to learn task-specific patterns
Epochs	5	Sufficient for transfer learning — adapting existing knowledge, not training from scratch
Batch Size	16	Fits within the 16 GB VRAM of the free Colab T4 GPU
Warmup Ratio	10%	Gradually increases learning rate at the start to prevent destabilizing weight updates
Weight Decay	0.01	L2 regularization to prevent overfitting
Mixed Precision	FP16	$\sim 2\times$ training speed with negligible accuracy impact
Best Model Metric	F1 Score	Balances precision and recall; trainer saves the checkpoint with highest F1

The ViTImageProcessor from Hugging Face handles the image-to-tensor conversion, including resizing to 224×224 and normalization using ImageNet statistics.

Training Results

The model was trained for 5 epochs, completing in approximately 15 minutes:

Epoch	Training Loss	Validation Loss	Accuracy	Precision	Recall	F1 Score
1	0.053199	0.023557	99.10%	99.10%	99.10%	99.10%
2	0.016501	0.051502	98.65%	97.37%	100.00%	98.69%
3	0.000173	0.031316	98.95%	97.94%	100.00%	98.96%
4	0.000072	0.018723	99.40%	98.81%	100.00%	99.40%
5	0.000065	0.019335	99.40%	98.81%	100.00%	99.40%

[Figure 7: table of the training output showing epoch-byepoch metrics]

Final Test Set Metric	Value
Accuracy	99.40%
Precision	98.81%
Recall	100.00%
F1 Score	99.40%

These results are very strong on the test set. However, it is important to note that real-world performance may differ from benchmark accuracy because the test set comes from the same distribution as the training data. When tested on images from generators not represented in the training data, or on images that have undergone post-processing (compression, resizing, filtering), the accuracy may decrease. This is a common challenge in AI detection and could be addressed in future work through expanded training data and additional data augmentation.

The trained model was pushed to Hugging Face Hub at [mohamed9679/ai-image-detector-v1](#) for deployment.

Evaluation Metrics Explained

For completeness, we briefly explain the metrics used:

- Accuracy: The percentage of all images correctly classified (both real and fake).
- Precision: Of all images predicted as fake, what percentage was actually fake. High precision means few false alarms.
- Recall: Of all actual fake images, what percentage did we catch. Our model achieved 100% recall, meaning it detected every fake image in the test set.
- F1 Score: The harmonic mean of precision and recall, providing a single balanced measure: $F1 = 2 \times (\text{Precision} \times \text{Recall}) / (\text{Precision} + \text{Recall})$

e. Ensemble Decision Methods

No single model is perfect. Ensemble methods combine the predictions of multiple models to produce a more robust and reliable decision. We implemented three different ensemble strategies across our system versions.

e.1 Simple Average Ensemble (V1)

The first version used the simplest possible ensemble: averaging the four model scores and applying a fixed threshold:

```
average_score=(resnet+siglip+sdxl+deepfake)/4.0
```

```
if average_score > 0.5:
```

```
    verdict="FAKE"
```

```
else:
```

```
    verdict="REAL"
```

This approach treats all models equally, regardless of their individual confidence levels. Its weakness is that a single highly confident but incorrect model can override three correct models.

e.2 Adaptive Confidence-Weighted Ensemble (V2)

The second version introduced a more sophisticated ensemble with three key improvements:

Confidence Weighting: Instead of simple averaging, each model's contribution is weighted by how confident it is. A model that outputs 0.95 (very confident) has more influence than one that outputs 0.52 (barely deciding):


$$w_i = |s_i - 0.5| + \epsilon$$

$$\text{weighted_avg} = \frac{\sum_i s_i \times w_i}{\sum_i w_i}$$

Adaptive Threshold: The decision threshold adjusts based on model agreement:

Condition	Threshold	Reasoning
$\geq 75\%$ models agree + low spread	0.45	Strong consensus — trust it even at slightly lower scores
$\geq 60\%$ models agree	0.50	Standard threshold
Models are split	0.55	Require stronger evidence to declare FAKE

Uncertainty Detection: When more than 60% of the models are individually uncertain (scores between 0.35 and 0.65) and there is high disagreement, the system outputs "UNCERTAIN" instead of forcing a binary choice.



[Figure 8 : Adaptive ensemble decision flow — threshold adjusts based on model agreement]

e.3 Fixed Weighted Ensemble (V3)

With a strong primary model (the fine-tuned ViT), we simplified the ensemble to use fixed weights. The fine-tuned model serves as the primary decision-maker while the backup models and noise forensics add robustness:

Engine	Weight	Role
Fine-tuned ViT	50%	Primary detector — trained for the specific task
SigLIP	15%	Backup — catches semantic-level inconsistencies
SMOGY	15%	Backup — complements with different training data
Noise Forensics	20%	Physics-based signal — catches non-ML artifacts

The weighted score is computed as:

$$\text{score} = 0.50 \times s_{\text{ViT}} + 0.15 \times s_{\text{SigLIP}} + 0.15 \times s_{\text{SMOGY}} + 0.20 \times s_{\text{noise}}$$

If the score exceeds 0.5, the verdict is "FAKE"; otherwise, it is "REAL."


The rationale for keeping backup models despite having a strong primary model is diversity: the noise forensic engine provides a fundamentally different type of signal (physics-based vs. machine learning-based), and it can detect anomalies that no neural network would catch. Similarly, the backup models were trained on different data distributions, providing complementary coverage.

f. Multi-Scale Inference (Test-Time Augmentation)

A limitation of standard image classification is that the model sees the image at only one scale and one orientation. Some AI generation artifacts are visible only at certain scales or become apparent when the image is slightly modified.

Test-Time Augmentation (TTA) addresses this by running the model on multiple views of the same image and averaging the results. We implemented this in V2 and V3 with different levels of sophistication.

V2: Four-View Multi-Scale



[Figure 9 : Multi-scale inference pipeline — four views are analyzed independently and averaged]

View	Transformation	What It Catches
Original	None	Baseline prediction
Horizontal Flip	Mirror image	Left/right asymmetry artifacts
Center Crop (80%)	Crop center region and resize back	Fine details that may be lost at full resolution
CLAHE Enhanced (V2 only)	Adaptive contrast enhancement	Low-contrast artifacts invisible in the original

V3: Three-View TTA

In the third version with the fine-tuned model, we simplified TTA to three views (original, flip, center crop), removing the CLAHE view since the fine-tuned model is already strong enough without it.

The multi-scale inference adds approximately 3-4× computation time per model, but significantly reduces the variance of predictions — a model that might score 0.48 (just below threshold) on the original image might score 0.55, 0.61, and 0.52 on the other views, making the average decisively above threshold.

g. Summary

We designed and implemented a multi-engine AI image detection system that evolved through three iterations, each addressing specific limitations of the previous version.

Key contributions of our implementation:

- 1. **Multi-modal analysis:** We combined visual forensic methods (FFT, ELA), physics-based analysis (noise forensics), and deep learning classification into a unified system. This diversity of approaches provides robustness — even if one method fails, others can compensate.

2. **Noise pattern forensics:** We implemented a novel noise analysis engine that extracts and analyzes four statistical features of image noise (variance, spatial correlation, channel consistency, and entropy). This physics-based approach requires no training data and catches artifacts that machine learning models may miss.
3. **Fine-tuned ViT model:** Rather than relying entirely on pre-trained models from other researchers, we fine-tuned our own Vision Transformer on a modern dataset. The fine-tuned model achieved 99.40% accuracy and 99.40% F1 score on the test set.

Limitations and future work:

While the fine-tuned model achieved excellent results on the benchmark test set, real-world testing with images from generators not represented in the training data showed room for improvement. Future work could include:

- Training on a larger and more diverse dataset covering more AI generators
 - Implementing periodic retraining as new generators emerge
- Adding data augmentation during training (compression, resizing, noise injection) to improve robustness to post-processing

Exploring larger ViT variants (ViT-Large, ViT-Huge) for potentially higher accuracy

3.3 Deployment and Interface Design

a. Employment of Vision Transformer (ViT)

The project utilizes **Hugging Face Spaces** as the primary hosting environment.

- **Containerization:** Spaces provide a specialized infrastructure that manages the hardware requirements (RAM and CPU/GPU) needed to run Transformer models efficiently.
- **Seamless Integration:** By connecting the GitHub repository directly to the Space, a "Continuous Deployment" (CD) pipeline is established. Any update to the app.py or model weights on GitHub is automatically built and deployed on the cloud.

a.2 Gradio Framework Architecture

To build the User Interface (UI), the **Gradio** library was selected for its efficiency in machine learning applications.

- **Rapid Prototyping:** Gradio allows for the creation of an interactive interface with minimal code, focusing on the predict() function that processes the user's image.
- **Input-Output Mapping:** The interface is designed to take a "PIL Image" as input and return a "Label" as output, which provides the most intuitive experience for a non-technical user.

a.3 User Interaction Flow and Experience (UX)

The user journey is designed to be straightforward and responsive:

- **Image Upload:** Users can drag and drop or upload an image in various formats (JPG, PNG, WebP).
- **Processing:** Once uploaded, the image is sent to the inference pipeline where the ViT model performs its analysis.

- **Real-time Feedback:** The system provides a "Probability Bar" showing the confidence level for both "Real" and "AI-Generated" classes, allowing users to see the model's certainty.

a.4 Frontend Design and Descriptive Elements

- The interface is not just a functional tool but also an educational one.
- **Title and Description:** The `gr.Interface` includes a clear title ("AI Image Detector") and a description that guides the user on how the system works.
- **Visual Clarity:** Using the `num_top_classes=2` parameter in the output ensures that the interface remains clean, showing only the most relevant results without cluttering the screen with raw data.

a.5 Backend Server Configuration

To ensure the application remains stable and accessible, specific server configurations were implemented in the `app.py` script:

- **Network Binding:** The command `demo.launch(server_name="0.0.0.0", server_port=7860)` is crucial for cloud deployment.
- **Accessibility:** Setting the server name to `0.0.0.0` allows the application to listen on all available network interfaces, which is a requirement for containerized hosting environments like Hugging Face or Docker.

a.6 Security and Model Persistence

Managing the model's large file size and security was a key consideration during deployment.

- **Weight Management:** The model weights are stored and called via the Hugging Face pipeline API, which optimizes the loading time and ensures that the Space doesn't crash due to memory overflows.
- **Authentication:** The use of scripts like `check_my_token.py` (as seen in the repository) ensures that the communication between the local environment and the cloud repository is authenticated and secure.

b Employment of (Multi-Engine Ensemble) Model

b.1 Streamlit Version (V1)

The first version was built with Streamlit, a Python framework for rapid prototyping of data applications. Streamlit handles the UI automatically from Python code — no HTML, CSS, or JavaScript is needed.

The application structure was modular:

```
webapp/
├── app.py          #MainStreamlitapplication
├── src/
│   ├── preprocessing.py #Imagepreparationpipeline
│   ├── analyzers.py     #FFT,ELA,andMLmodelfunctions
│   └── ensemble.py      #Verdictcalculation
└── requirements.txt
```

The Streamlit version provided: - File upload for images (JPG, PNG, WEBP) - Side-by-side display of FFT and ELA results with interpretation guides - Individual model confidence scores as metric widgets - A “Final Judge” section showing the ensemble verdict

Model caching was handled by Streamlit’s built-in `@st.cache_resource` decorator, which loads each model only once and reuses it across sessions.

b.2Gradio + FastAPI Version (V2 and V3)

We migrated to Gradio with FastAPI for the second and third versions for the following reasons:

Requirement	Streamlit	Gradio + FastAPI
REST API for mobile app	✗ Not built-in	✓ FastAPI provides /analyze endpoint
Hugging Face Spaces hosting	⚠ Limited support	✓ Native support
Custom UI themes	⚠ Limited	✓ Full CSS customization
API documentation	✗ None	✓ Auto-generated Swagger docs

The combined architecture mounts the Gradio UI on the root path ("/") and exposes a RESTful API endpoint at "/analyze":

```
fastapi_app=FastAPI(title="AllImageDetectorAPI")

@fastapi_app.post("/analyze")
async def analyze_endpoint(file:UploadFile=File(...)):
    content=await file.read()
    pil_img=Image.open(io.BytesIO(content)).convert("RGB")
    result=run_full_analysis(pil_img)
    api_result={k:vfork,vinresult.items()ifnotk.startswith("_")}
    return JSONResponse(content=api_result)

app=gr.mount_gradio_app(fastapi_app,demo,path="/")
```

The REST API accepts an image as multipart form data and returns a JSON response containing the verdict, confidence score, and individual model scores. This enables integration with mobile applications (e.g., an Android app) that can send images programmatically and receive structured results.

API Response Format:

```
{  
  "verdict": "FAKE", "confidence": 87.45,  
  "agreement": "Strongconsensus(5fake/2real)", "scores": {  
    "resnet": 92.3,  
    "siglip": 88.1,  
    "sdxl": 85.0,  
    "deepfake": 84.5,  
    "wvolf": 97.2,  
    "smogy": 95.8,  
    "noise": 72.3  
  }  
}
```

b.3UI Design (V3)

The third version featured a premium UI design with custom CSS styling:

- Dark gradient background with purple accent colors
- Custom Inter font from Google Fonts
- Engine cards showing each model's weight and role at a glance
- Glassmorphism-style containers with backdrop blur effects
- Animated hover effects on buttons
- Color-coded verdict display (red for FAKE, green for REAL)

These design choices were made to create a professional, modern-looking application appropriate for demonstrating the project's capabilities

b.4 Hugging Face Spaces

The application was deployed on Hugging Face Spaces, a free hosting platform for machine learning demos. Hugging Face Spaces provides:

- Free GPU/CPU runtime for inference
- Automatic dependency installation from requirements.txt
- HTTPS endpoint with a public URL
- Git-based deployment

The deployment configuration is specified in the README.md metadata:

```
title:AllImageDetector
sdk:gradio
app_file:app.py
```

For the fine-tuned model version (V3), a Dockerfile was created to ensure consistent deployment. Docker containerization guarantees that the application runs identically regardless of the hosting environment, eliminating environment-specific issues.

2.5 Docker Containerization

For the fine-tuned model version (V3), we created a Dockerfile to ensure consistent deployment:

```
FROMpython:3.10-slim
WORKDIR/app
COPYrequirements.txt.
RUNpipinstall--no-cache-dir-rrequirements.txt
COPY..
CMD["uvicorn","app:app","--host","0.0.0.0","--port","7860"]
```

Docker containerization guarantees that the application runs identically regardless of the hosting environment, eliminating "it works on my machine" issues.

b.6CI/CD with GitHub Actions

We set up a GitHub Actions workflow to automatically sync the repository to Hugging Face Spaces on every push to the main branch:

```
name:SynctoHuggingFace
on:
  push:
    branches:[main]jobs:
  sync:
    runs-on:ubuntu-latest
    steps:
      - uses:actions/checkout@v3
      - name:PushtoHFSpace
        run:|
          gitpush
          https://USER:${{secrets.HF_TOKEN}}@huggingface.co/spaces/...main--force
```

This automation ensures that any code changes pushed to GitHub are automatically deployed to the live application without manual intervention.

c. Implementing the Models in an Android App

c.1 Architectural Design: Model-View-ViewModel (MVVM)

The application is built upon the MVVM architectural pattern. The View Layer (Fragments) handles the user interface and observing data changes with no business logic. The ViewModel Layer (UploadViewModel) acts as a mediator between View and Model, surviving configuration changes and holding application state. The Model Layer handles data operations including image processing and network communication with the Hugging Face inference API.

c.2Concurrency and Asynchronous Processing: Kotlin Coroutines

To maintain a responsive UI, the system utilizes Kotlin Coroutines for network requests and heavy image processing. Using `viewModelScope`, the application launches asynchronous tasks that perform API calls in the background, with results safely posted back to the UI thread via `LiveData`.

c.3Network Communication and Remote Inference Integration

The application targets a specific REST API endpoint hosted on Hugging Face Spaces. A modern, asynchronous HTTP client (Ktor) is configured to handle POST requests using the CIO engine for high-performance, non-blocking network I/O.

c.4 Security and Authentication: API Access Tokens

Access to the Hugging Face API is secured using Bearer Token Authentication. The token is managed via the `local.properties` file and injected through the `BuildConfig` class, preventing sensitive credentials from being exposed in version control systems.

c.5 Operational Workflow

The operational workflow consists of: (1) Connectivity Validation; (2) Image Pre-processing; (3) Data Transmission — the processed image is encoded into Base64, bundled into JSON, and transmitted with the Authorization Header; and (4) Inference and Visualization — the remote AI model analyzes the image, calculates the AI Confidence percentage, and navigates the user to the result screen.

c.6 Summary

AlBuster represents a robust integration of modern Android development practices. By combining the MVVM pattern with Coroutines and Ktor, the application achieves a high standard of performance and reliability, providing a scalable foundation for future enhancements in the field of AI detection.

C H A P T E R 4

Comparative Analysis of AI Image Detection Models

This chapter presents a comparative study of two distinct approaches to AI-generated image detection. Both systems share a common foundation — the Vision Transformer (ViT) architecture — yet differ substantially in their design philosophy, forensic methodology, and deployment strategy. The first model adopts a focused, single-engine approach centered on a fine-tuned ViT, while the second model constructs a multi-layered ensemble that combines deep learning with classical signal processing and physics-based forensics.

4.1 Model 1 — The Standalone Vision Transformer (ViT)

Overview

The first model is built around a single, purpose-fine-tuned Vision Transformer. Rather than relying on an ensemble of independent detectors, this approach concentrates its detection capability into one highly optimized network. The model was pre-trained on ImageNet-21k to develop broad visual understanding, then fine-tuned specifically on AI detection datasets including CIFAKE, Dima806/AI_vs_Real, and real-world images sourced through the CapCheck repository.

How It Works

The model processes input images by dividing them into fixed 16×16 pixel patches. Each patch is flattened into a vector, projected into a latent embedding space, and combined with positional encodings to preserve spatial structure. A special [CLS] classification token is prepended to the sequence and passed through a stack of Transformer encoder layers. At each layer, the Multi-Head Self-Attention mechanism computes relationships between every patch simultaneously, allowing the model to detect global inconsistencies that local convolutional filters would miss. A lowered classification threshold of 0.08 makes the detector highly sensitive to weak generative signals, boosting recall at the cost of a marginal increase in false positives.

Strengths

The primary strength lies in global contextual awareness — by attending to all patches simultaneously, it can identify structural artifacts spanning the entire image. The architecture benefits from reduced inductive bias, making it more adaptable to new generative engines as training data is updated. The single-model design is computationally lean, offering fast and predictable inference (45ms) suitable for real-time applications.

Limitations

The model's main limitation is its reliance on a single modality of analysis, producing no human-interpretable forensic output. The aggressive 0.08 threshold increases the risk of false positives on real photographs with unusual textures or synthetic-looking lighting. Performance may degrade on images from generative engines not represented in the fine-tuning dataset.

4.2 Model 2 — The Multi-Engine Ensemble Detection System

Overview

The second model combines multiple independent analysis engines — each examining the image from a different perspective — and aggregates their outputs into a single weighted verdict. The system evolved across three development iterations (V1 through V3), with the final version (V3) employing a fine-tuned ViT as the primary detector, augmented by two backup deep learning models and a physics-based noise forensics engine.

How It Works

Every uploaded image passes through a preprocessing pipeline that strips EXIF metadata, generates grayscale and JPEG-compressed variants, and optionally applies CLAHE contrast enhancement. The fine-tuned ViT (trained on the Parveshiiii/Al-vs-Real dataset) serves as the primary classifier contributing 50% of the ensemble weight. A SigLIP model and the SMOGY detector each

contribute 15%. The noise forensics engine contributes 20% and operates entirely without machine learning — extracting statistical properties of the image's noise residual. FFT and ELA provide additional human-interpretable outputs, and Test-Time Augmentation further improves robustness.

Strengths

The multi-engine design provides layered redundancy. The physics-based noise forensics engine can detect anomalies in images from generators never seen during training, addressing a core vulnerability of purely ML-based approaches. The system offers superior transparency through FFT and ELA visual outputs. The fine-tuned ViT achieved 99.40% accuracy and 100% recall on its test set.

Limitations

The multi-engine architecture carries significantly higher computational cost. TTA multiplies inference time by three to four times. Deployment complexity is greater, requiring coordination between multiple model checkpoints and a REST API layer. Real-world performance may be lower for images that have undergone post-processing such as compression, resizing, or noise injection.

4.3 Side-by-Side Comparison

Criterion	Model 1 — ViT Standalone	Model 2 — Multi-Engine Ensemble
Architecture	Single Vision Transformer (ViT-base-patch16-224)	Multi-engine ensemble: ViT fine-tuned + SigLIP + SMOGY + Noise Forensics
Core Mechanism	Multi-Head Self-Attention across image patches	Adaptive ensemble combining deep learning and signal processing
Input Processing	Image divided into 16×16 patches → linear	Preprocessing pipeline: metadata stripping, CLAHE,

Criterion	Model 1 — ViT Standalone	Model 2 — Multi-Engine Ensemble
	projection → positional embedding	JPEG re-compression for ELA
Training Approach	Transfer learning: pre-trained on ImageNet-21k, fine-tuned on CIFAKE + Dima806	Fine-tuned ViT (Parveshiiii dataset) + pre-trained backup models, no joint training
Dataset Used	CIFAKE, Dima806/Al_vs_Real, CapCheck (in-the-wild images)	Parveshiiii/Al-vs-Real (~10,000 images, 80/20 split)
Accuracy	94.2%	99.40% (test set)
F1-Score	94.2%	99.40%
Precision	Not separately reported	98.81%
Recall	Not separately reported	100.00%
Inference Speed	~45ms per image	Higher latency due to multi-engine pipeline (3–4× TTA overhead)
Threshold Strategy	Custom threshold of 0.08 (lowered from default 0.5) to improve recall	Fixed weighted ensemble threshold of 0.5; V2 uses adaptive threshold
Forensic Analysis	None — relies solely on learned attention patterns	FFT (frequency artifacts), ELA (compression inconsistency), Noise Pattern Analysis (PRNU)
Explainability	Moderate — attention maps can visualize patch	Higher — FFT and ELA outputs are human-

Criterion	Model 1 — ViT Standalone	Model 2 — Multi-Engine Ensemble
	importance	interpretable visual evidence
Ensemble Method	Single model — no ensemble	V1: simple average; V2: confidence-weighted adaptive; V3: fixed weights
Robustness to Novel Generators	Moderate — may degrade on unseen generators	Noise forensics provides physics-based detection independent of generator type
Deployment Platform	Hugging Face Spaces + Gradio UI	Gradio + FastAPI with REST API (supports web and mobile)
Scalability	Excellent at scale; ViT outperforms CNN as dataset grows	Computationally heavier due to multiple parallel engines and TTA
Data Augmentation	Balanced sampling; error handling for corrupted files	CLAHE enhancement, multi-scale TTA (3–4 views per image)
Main Strength	High global context understanding; detects structural and spectral artifacts holistically	Multi-modal redundancy; physics-based noise analysis catches what ML models miss
Main Limitation	Single modality; no forensic visual output; threshold tuning required	Higher complexity; real-world accuracy may drop on post-processed or compressed images

4.4 Discussion

The two models represent complementary points on the spectrum between simplicity and comprehensiveness. Model 1 demonstrates that a well-fine-tuned single ViT, with a carefully tuned decision threshold, can achieve strong and efficient detection suitable for deployment in resource-constrained or latency-sensitive environments. Its 94.2% accuracy is competitive, and its 45ms inference time makes it viable for real-time applications.

Model 2 trades computational efficiency for robustness and interpretability. Its 99.40% test accuracy is meaningfully higher, and the inclusion of physics-based noise forensics provides a detection capability that no neural network alone can replicate. This makes it particularly well-suited for high-stakes forensic contexts where both accuracy and explainability are required, such as journalism verification, legal evidence review, or content moderation platforms.

It is worth noting, however, that the two models were evaluated on different datasets, which makes direct numerical comparison imprecise. Model 1 was tested on CIFAKE and Dima806, while Model 2 used the Parveshiiii dataset. The 5-percentage-point accuracy gap may therefore reflect differences in dataset difficulty as much as architectural superiority. Both models share the potential limitation of reduced generalization to images from generative engines not covered in their respective training sets, and both would benefit from ongoing retraining as the generative AI landscape evolves.

In summary, the choice between these two architectures depends on the deployment context: Model 1 offers speed and simplicity, while Model 2 offers depth, redundancy, and interpretability.

C H A P T E R 5

Conclusion and Future Work

5.1 Summary of Work

This research set out to address one of the most pressing challenges in the modern information ecosystem: the reliable and scalable detection of AI-generated imagery. As generative AI tools become increasingly accessible and their outputs increasingly indistinguishable from authentic photographs, the need for robust, explainable, and deployable detection systems has become critical.

Over the course of this project, the team designed, implemented, and evaluated two fundamentally different detection architectures — a standalone fine-tuned Vision Transformer (ViT) and a multi-engine ensemble system — and deployed them through both a web application and an Android mobile interface. The project encompassed the full pipeline from theoretical justification and data preparation, through model training and evaluation, to cloud deployment and mobile integration.

The first model (3.1) demonstrated that a single, carefully optimized ViT, pre-trained on ImageNet-21k and fine-tuned on specialized AI detection datasets, can achieve strong performance (94.2% accuracy, ~45ms inference) with a low-complexity deployment footprint. A key implementation insight was the reduction of the classification threshold from 0.5 to 0.08, which significantly improved recall by making the model sensitive to weak generative signals. This model is well-suited for real-time, resource-constrained environments.

The second model (3.2) advanced the state of detection by integrating four parallel analysis engines: a fine-tuned ViT as the primary classifier (50% weight), two backup deep learning models providing complementary coverage, and a novel physics-based noise forensics engine (20% weight). The noise forensics engine is a particularly noteworthy contribution — by analyzing four statistical properties of the image noise residual (variance, spatial correlation, channel consistency, and Shannon entropy), it detects the unnaturally structured noise patterns characteristic of AI-generated images without relying on any training data. The ensemble achieved 99.40% accuracy and 100% recall on its test set,

and its three-version evolution (V1 through V3) demonstrated a disciplined engineering approach to progressively improving robustness, interpretability, and deployment quality.

Finally, the Android application (Section 3.3.3) translated these detection capabilities into a mobile experience, using an MVVM architecture, Kotlin Coroutines for asynchronous processing, and Ktor for secure API communication with the Hugging Face-hosted inference engine. This step completes the project's goal of producing a detection tool that is not only academically rigorous but practically deployable by real end-users.

5.2 Key Findings

Finding	Details
ViT outperforms CNN for global artifact detection	ViT's Multi-Head Self-Attention provided a 3–5% accuracy boost over standard CNN baselines (ResNet, VGGNet, DenseNet) on the same dataset, by capturing long-range structural inconsistencies that local convolutional filters miss
Physics-based analysis complements ML	The noise forensics engine detects artifacts invisible to neural networks, particularly the structured, low-entropy, channel-correlated noise patterns that AI generators produce — and does so without requiring any training data
Ensemble diversity improves robustness	No single model is universally best; combining models trained on different data distributions with a physics-based engine creates a detection system that is more robust to novel generators and post-processing attacks
Threshold tuning is critical	Reducing the classification threshold from 0.5 to 0.08 in Model 1 significantly improved recall without unacceptable increases in false positives,

Finding	Details
	highlighting that default thresholds are not always optimal for high-stakes detection tasks
Test-Time Augmentation reduces prediction variance	Running models on multiple views (original, flipped, center-cropped) and averaging predictions significantly reduces the chance of borderline misclassifications near the decision boundary
Dataset distribution affects benchmark comparability	Models trained and tested on different datasets cannot be directly compared numerically; the 5-percentage-point accuracy gap between models may reflect dataset difficulty as much as architectural superiority

5.3 Limitations

Despite the strong results achieved, several important limitations should be acknowledged:

Distribution gap: Both models were evaluated on test sets drawn from the same distribution as their training data. When tested on images from generators not represented in the training set, or on images that have undergone post-processing (JPEG compression, resizing, noise injection, social media re-encoding), accuracy decreases. This is a fundamental challenge in AI detection because generative models continuously improve and new generators emerge.

Dataset incomparability: The two models were trained and evaluated on different datasets (CIFAKE/Dima806 for Model 1; Parveshiiii for Model 2), making direct numerical comparison of accuracy figures imprecise.

Computational cost: The multi-engine system with Test-Time Augmentation imposes a 3–4× latency overhead compared to a single-model approach. This may limit deployment in real-time or mobile-native contexts.

Threshold sensitivity: The aggressive 0.08 threshold in Model 1, while improving recall, increases the rate of false positives on real photographs with unusual visual characteristics. Further calibration studies on diverse real-world data would be needed before production deployment.

Android inference delegation: The Android application delegates inference to a remote API rather than running the model on-device, introducing latency and a dependency on network connectivity and server availability.

5.4 Future Work

The following directions represent the most promising avenues for extending and improving this work:

Direction	Description
Expanded and continuously updated training data	Training on a larger, more diverse dataset covering more AI generators (including emerging models) and implementing periodic retraining pipelines to keep pace with new generation techniques
Data augmentation for robustness	Systematically applying compression, resizing, cropping, noise injection, and color jittering during training to improve the models' resilience to post-processing operations that are commonly applied to shared images
Larger ViT variants	Exploring ViT-Large and ViT-Huge architectures, which offer more parameters and potentially higher representational capacity for detecting subtle generative artifacts
On-device mobile inference	Replacing the remote API call in the Android app with an on-device model (e.g., a quantized MobileViT) to improve latency, offline capability, and user privacy

Direction	Description
C2PA and SynthID integration	Combining detection-based approaches with provenance-based standards (C2PA) and watermarking systems (SynthID) to create a layered authenticity verification system that is both reactive and proactive
Explainability enhancements	Developing richer visualization tools — such as attention map overlays, Grad-CAM heatmaps, and noise residual annotations — to provide users with more granular and interpretable evidence for each verdict
Cross-generator generalization benchmarking	Establishing a standardized evaluation protocol that tests models on images from generators they have not seen during training, providing a more realistic measure of real-world robustness
Multimodal detection extension	Extending the detection framework beyond static images to cover AI-generated video (deepfakes), audio, and text, which increasingly coexist and are combined in synthetic media attacks

5.5 Closing Remarks

The landscape of AI-generated media is evolving at a pace that makes static detection solutions inadequate. Any model trained today will face new generators tomorrow. What this project establishes, however, is a principled and extensible detection framework: one that combines the global pattern recognition of Vision Transformers, the interpretable evidence of visual forensics, and the generalization power of physics-based signal analysis.

The AlBuster project demonstrates that effective AI detection does not require choosing between accuracy and interpretability, or between performance and deployability. By thoughtfully combining multiple complementary approaches


— and by designing the system to be iteratively improved — it is possible to build a detection tool that is both rigorous enough to satisfy researchers and practical enough to serve real users.

As generative AI becomes ever more embedded in daily life, the importance of trustworthy, accessible, and explainable detection tools will only grow. This project contributes a concrete step toward that goal, and lays the groundwork for the next generation of AI authenticity verification systems.


C H A P T E R 6

Diagrams


Use Case Diagram



Class Diagram



Activity Diagram



References

References

References of Chapter 1 & 2

- [1] Washington University in St. Louis, Information Security. "The Dangers of AI Art and Deepfakes." Last Accessed: March 2026.
- [2] MDPI. "Applied Sciences," vol. 15, no. 13, article 7421. Last Accessed: March 2026.
- [3] C2PA Viewer. "What is C2PA?" Last Accessed: March 2026.
- [4] G. Renish Rai, "SynthID Explained: A Technical Deep Dive into DeepMind's Invisible Watermarking System," dev.to. Last Accessed: March 2026.
- [5] Google DeepMind. "SynthID." Last Accessed: March 2026.
- [6] The Moonlight. "Methods and Trends in Detecting Generated Images: A Comprehensive Review." Last Accessed: March 2026.

Resources of Chapter 3

1. AI-Generated Image Detection

1.1 General

- [7] Z. Sha, Z. Li, N. Yu, and Y. Zhang, "DE-FAKE: Detection and Attribution of Fake Images Generated by Text-to-Image Generation Models," in *Proc. 2023 ACM SIGSAC Conf. Computer and Communications Security (CCS '23)*, pp. 3418–3432, 2023. DOI: 10.1145/3576915.3616588.
- [8] S.-Y. Wang, O. Wang, R. Zhang, A. Owens, and A. A. Efros, "CNN-Generated Images Are Surprisingly Easy to Spot...for Now," in *Proc. IEEE/CVF Conf. Computer Vision and Pattern Recognition (CVPR 2020)*, pp. 8695–8704, 2020. DOI: 10.1109/CVPR42600.2020.00872.
- [9] R. Corvi, D. Cozzolino, G. Zingarini, G. Poggi, K. Nagano, and L. Verdoliva, "On the Detection of Synthetic Images Generated by Diffusion Models," in *IEEE Int. Conf. Acoustics, Speech, and Signal Processing (ICASSP 2023)*, pp. 1–5, 2023. DOI: 10.1109/ICASSP49357.2023.10095167.

1.2 GAN-Generated Image Detection

- [10] F. Marra, D. Gragnaniello, L. Verdoliva, and G. Poggi, "Do GANs Leave Artificial Fingerprints?" in *IEEE Conf. Multimedia Information Processing and Retrieval (MIPR 2019)*, pp. 506–511, 2019. DOI: 10.1109/MIPR.2019.00103.
- [11] N. Yu, L. S. Davis, and M. Fritz, "Attributing Fake Images to GANs: Learning and Analyzing GAN Fingerprints," in *Proc. IEEE/CVF Int. Conf. Computer Vision (ICCV 2019)*, pp. 7556–7566, 2019. DOI: 10.1109/ICCV.2019.00765.

2. Visual Forensic Techniques

2.1 Fast Fourier Transform (FFT) for Image Forensics

- [12] J. W. Cooley and J. W. Tukey, "An Algorithm for the Machine Calculation of Complex Fourier Series," *Mathematics of Computation*, vol. 19, no. 90, pp. 297–301, 1965. DOI: 10.2307/2003354.
- [13] T. Dzanic, K. Shah, and F. Witherden, "Fourier Spectrum Discrepancies in Deep Network Generated Images," in *Advances in Neural Information Processing Systems (NeurIPS 2020)*, vol. 33, pp. 3022–3032, 2020.
- [14] J. Frank, T. Eisenhofer, L. Schönherr, A. Fischer, D. Kolossa, and T. Holz, "Leveraging Frequency Analysis for Deep Fake Image Recognition," in *Proc. 37th Int. Conf. Machine Learning (ICML 2020)*, PMLR 119, pp. 3247–3258, 2020.

2.2 Error Level Analysis (ELA)

- [15] N. Krawetz, "A Picture's Worth... Digital Image Analysis and Forensics," in *Black Hat Briefings USA 2007*, 2007. Last Accessed: March 2026.
- [16] G. K. Wallace, "The JPEG Still Picture Compression Standard," *IEEE Trans. Consumer Electronics*, vol. 38, no. 1, pp. xviii–xxxiv, 1992. DOI: 10.1109/30.125072.

2.3 Noise Pattern / PRNU Forensics

- [17] J. Lukáš, J. Fridrich, and M. Goljan, "Digital Camera Identification from Sensor Pattern Noise," *IEEE Trans. Information Forensics and Security*, vol. 1, no. 2, pp. 205–214, 2006. DOI: 10.1109/TIFS.2006.873602.
- [18] M. Chen, J. Fridrich, M. Goljan, and J. Lukáš, "Determining Image Origin and Integrity Using Sensor Noise," *IEEE Trans. Information Forensics and Security*, vol. 3, no. 1, pp. 74–90, 2008. DOI: 10.1109/TIFS.2007.916285.

2.4 Median Filter Denoising

- [19] T. S. Huang, G. J. Yang, and G. Y. Tang, "A Fast Two-Dimensional Median Filtering Algorithm," *IEEE Trans. Acoustics, Speech, and Signal Processing*, vol. 27, no. 1, pp. 13–18, 1979. DOI: 10.1109/TASSP.1979.1163188.

2.5 CLAHE (Contrast-Limited Adaptive Histogram Equalization)

- [20] K. Zuiderveld, "Contrast Limited Adaptive Histogram Equalization," in *Graphics Gems IV*, Academic Press, pp. 474–485, 1994. ISBN: 0-12-336155-9.
- [21] S. M. Pizer et al., "Adaptive Histogram Equalization and Its Variations," *Computer Vision, Graphics, and Image Processing*, vol. 39, no. 3, pp. 355–368, 1987. DOI: 10.1016/S0734-189X(87)80186-X.

2.6 Shannon Entropy

- [22] C. E. Shannon, "A Mathematical Theory of Communication," *The Bell System Technical Journal*, vol. 27, pp. 379–423, 623–656, 1948. DOI: 10.1002/j.1538-7305.1948.tb01338.x.

3. Deep Learning Architectures

3.1 Vision Transformer (ViT)

- [23] A. Dosovitskiy et al., "An Image Is Worth 16x16 Words: Transformers for Image Recognition at Scale," in *Int. Conf. Learning Representations (ICLR 2021)*, 2021.
- [24] A. Steiner, A. Kolesnikov, X. Zhai, R. Wightman, J. Uszkoreit, and L. Beyer, "How to Train Your ViT? Data, Augmentation, and Regularization in Vision Transformers," *Transactions on Machine Learning Research (TMLR)*, 2022.

3.2 ResNet

- [25] K. He, X. Zhang, S. Ren, and J. Sun, "Deep Residual Learning for Image Recognition," in *Proc. IEEE Conf. Computer Vision and Pattern Recognition (CVPR 2016)*, pp. 770–778, 2016. DOI: 10.1109/CVPR.2016.90.

3.3 SigLIP

- [26] X. Zhai, B. Mustafa, A. Kolesnikov, and L. Beyer, "Sigmoid Loss for Language Image Pre-Training," in *Proc. IEEE/CVF Int. Conf. Computer Vision (ICCV 2023)*, pp. 11975–11986, 2023. DOI: 10.1109/ICCV51070.2023.01100.

3.4 Transformer Architecture

- [27] A. Vaswani et al., "Attention Is All You Need," in *Advances in Neural Information Processing Systems (NeurIPS 2017)*, vol. 30, pp. 5998–6008, 2017.

4. Transfer Learning and Fine-Tuning

- [28] S. J. Pan and Q. Yang, "A Survey on Transfer Learning," *IEEE Trans. Knowledge and Data Engineering*, vol. 22, no. 10, pp. 1345–1359, 2010. DOI: 10.1109/TKDE.2009.191.
- [29] J. Deng, W. Dong, R. Socher, L.-J. Li, K. Li, and L. Fei-Fei, "ImageNet: A Large-Scale Hierarchical Image Database," in *IEEE Conf. Computer Vision and Pattern Recognition (CVPR 2009)*, pp. 248–255, 2009. DOI: 10.1109/CVPR.2009.5206848.

5. Ensemble Methods

- [30] T. G. Dietterich, "Ensemble Methods in Machine Learning," in *Multiple Classifier Systems (MCS 2000)*, Lecture Notes in Computer Science, vol. 1857, pp. 1–15, Springer, 2000. DOI: 10.1007/3-540-45014-9_1.
- [31] J. Kittler, M. Hatef, R. P. W. Duin, and J. Matas, "On Combining Classifiers," *IEEE Trans.*

Pattern Analysis and Machine Intelligence, vol. 20, no. 3, pp. 226–239, 1998. DOI: 10.1109/34.667881.

6. Test-Time Augmentation (TTA)

- [32] A. Krizhevsky, I. Sutskever, and G. E. Hinton, "ImageNet Classification with Deep Convolutional Neural Networks," in *Advances in Neural Information Processing Systems (NeurIPS 2012)*, vol. 25, pp. 1097–1105, 2012.
- [33] D. Shanmugam, D. Blalock, G. Balakrishnan, and J. Guttag, "Better Aggregation in Test-Time Augmentation," in *Proc. IEEE/CVF Int. Conf. Computer Vision (ICCV 2021)*, pp. 1214–1223, 2021. DOI: 10.1109/ICCV48922.2021.00125.

7. Evaluation Metrics

- [34] C. J. Van Rijsbergen, *Information Retrieval*, 2nd ed. Butterworths, London, 1979. ISBN: 0-408-70929-4.
- [35] D. M. W. Powers, "Evaluation: From Precision, Recall and F-Measure to ROC, Informedness, Markedness and Correlation," *Journal of Machine Learning Technologies*, vol. 2, no. 1, pp. 37–63, 2011.

8. Dataset

- [36] Parveshiiii, "AI-vs-Real Dataset," Hugging Face Datasets. Last Accessed: March 2026.

9. Pre-Trained Models

- [37] umm-maybe, "AI-image-detector," Hugging Face Hub. Last Accessed: March 2026.
- [38] Ateeqq, "ai-vs-human-image-detector," Hugging Face Hub. Last Accessed: March 2026.
- [39] Organika, "sdxl-detector," Hugging Face Hub. Last Accessed: March 2026.
- [40] prithivMLmods, "Deep-Fake-Detector-v2-Model," Hugging Face Hub. Last Accessed: March 2026.
- [41] Wvolf, "ViT_Deepfake_Detection," Hugging Face Hub. Last Accessed: March 2026.
- [42] Smogy, "SMOGY-Ai-images-detector," Hugging Face Hub. Last Accessed: March 2026.
- [43] Google, "vit-base-patch16-224," Hugging Face Hub. Last Accessed: March 2026.
- [44] Mohamed9679, "ai-image-detector-v1 (Our Fine-Tuned Model)," Hugging Face Hub. Last Accessed: March 2026.

10. Frameworks and Libraries

10.1 Deep Learning and ML

- [45] A. Paszke et al., "PyTorch: An Imperative Style, High-Performance Deep Learning Library,"

in *NeurIPS 2019*, vol. 32, pp. 8024–8035, 2019.

- [46] T. Wolf et al., "Transformers: State-of-the-Art Natural Language Processing," in *EMNLP 2020 — System Demonstrations*, pp. 38–45, 2020.
- [47] Q. Lhoest et al., "Datasets: A Community Library for Natural Language Processing," in *EMNLP 2021 — System Demonstrations*, pp. 175–184, 2021.
- [48] Hugging Face. "Evaluate Library." Last Accessed: March 2026.
- [49] Hugging Face. "Accelerate Library." Last Accessed: March 2026.
- [50] F. Pedregosa et al., "Scikit-learn: Machine Learning in Python," *Journal of Machine Learning Research (JMLR)*, vol. 12, pp. 2825–2830, 2011.

10.2 Image Processing

- [51] Pillow (PIL) Contributors, "Pillow: Python Imaging Library (Fork)." Last Accessed: March 2026.
- [52] C. R. Harris et al., "Array Programming with NumPy," *Nature*, vol. 585, pp. 357–362, 2020. DOI: 10.1038/s41586-020-2649-2.
- [53] J. D. Hunter, "Matplotlib: A 2D Graphics Environment," *Computing in Science & Engineering*, vol. 9, no. 3, pp. 90–95, 2007. DOI: 10.1109/MCSE.2007.55.

10.3 Web Frameworks

- [54] Streamlit Inc. "Streamlit." Last Accessed: March 2026.
- [55] A. Abid, A. Abdalla, A. Abid, D. Khan, A. Alfozan, and J. Zou, "Gradio: Hassle-Free Sharing and Testing of ML Models in the Wild," in *ICML 2019 Workshop*, 2019.
- [56] S. Ramírez, "FastAPI." Last Accessed: March 2026.
- [57] Encode OSS. "Uvicorn." Last Accessed: March 2026.

10.4 Training Infrastructure

- [58] Google. "Google Colab." Last Accessed: March 2026.
- [59] Hugging Face. "Hugging Face Hub." Last Accessed: March 2026.
- [60] Hugging Face. "Hugging Face Spaces." Last Accessed: March 2026.

11. Deployment Technologies

- [61] Docker Inc. "Docker." Last Accessed: March 2026.
- [62] GitHub. "GitHub Actions." Last Accessed: March 2026.
- [63] Software Freedom Conservancy. "Git." Last Accessed: March 2026.

12. AI Image Generators Referenced

- [64] R. Rombach, A. Blattmann, D. Lorenz, P. Esser, and B. Ommer, "High-Resolution Image Synthesis with Latent Diffusion Models," in *Proc. IEEE/CVF Conf. Computer Vision and Pattern Recognition (CVPR 2022)*, pp. 10684–10695, 2022. DOI: 10.1109/CVPR52688.2022.01042.
- [65] OpenAI. "DALL-E 3." Last Accessed: March 2026.
- [66] Midjourney, Inc. "Midjourney." Last Accessed: March 2026.
- [67] Black Forest Labs. "Flux." Last Accessed: March 2026.
- [68] T. Karras, T. Aila, S. Laine, and J. Lehtinen, "Progressive Growing of GANs for Improved Quality, Stability, and Variation," in *Int. Conf. Learning Representations (ICLR 2018)*, 2018.
- [69] T. Karras, S. Laine, and T. Aila, "A Style-Based Generator Architecture for Generative Adversarial Networks," in *Proc. IEEE/CVF Conf. Computer Vision and Pattern Recognition (CVPR 2019)*, pp. 4401–4410, 2019.

13. Related Commercial Systems

- [70] SightEngine. "AI Image Detection API." Last Accessed: March 2026.

14. Design Resources

- [71] Google Fonts. "Inter Typeface." Last Accessed: March 2026.

15. Additional Key References

15.1 Generative Adversarial Networks (GANs)

- [72] I. J. Goodfellow et al., "Generative Adversarial Networks," in *Advances in Neural Information Processing Systems (NeurIPS 2014)*, vol. 27, pp. 2672–2680, 2014.

15.2 Diffusion Models

- [73] J. Ho, A. Jain, and P. Abbeel, "Denoising Diffusion Probabilistic Models," in *Advances in Neural Information Processing Systems (NeurIPS 2020)*, vol. 33, pp. 6840–6851, 2020.
- [74] R. Rombach et al., "High-Resolution Image Synthesis with Latent Diffusion Models," in *Proc. IEEE/CVF Conf. Computer Vision and Pattern Recognition (CVPR 2022)*, pp. 10684–10695, 2022. DOI: 10.1109/CVPR52688.2022.01042.

15.3 Mixed Precision Training

- [75] P. Micikevicius et al., "Mixed Precision Training," in *Int. Conf. Learning Representations (ICLR 2018)*, 2018.

15.4 Pearson Correlation Coefficient

- [76] K. Pearson, "Notes on Regression and Inheritance in the Case of Two Parents," *Proc. Royal*

Society of London, vol. 58, pp. 240–242, 1895. DOI: 10.1098/rspl.1895.0041.

Last updated: March 2026